

Agent + Context Sources — Findings Snapshot

Does feeding an AI agent more info + autonomy beat just following model_v4?

SETUP

- Harness owns all fills/accounting; agent emits only structured trades, one day at a time, no look-ahead. strict = follow the model's queue exactly (fixed reference). balanced = mostly follow. discretionary = use own judgment. Scored as return vs strict.

Sonnet 4.6 — full window 2026-01-05 → 2026-06-16

Reference strict = +69.84% (≈ model +69.61%) | SPY +9.40% · QQQ +18.25% — every agent cell trailed strict.

Setting (vs strict)	none	+news	+macro	+macro&news
balanced	-8.47%	-7.31%	-0.83%	-0.44%
discretionary	-51.55%	-51.31%	-42.41%	-32.87%

Key findings (deltas are 30-50pp — well above noise)

- Nothing beat strict. On a momentum bull tape the agent adds no alpha— only subtracts.
- Autonomy is destructive: discretion churned (100-160 trades) and bled -33 to -52pp vs strict.
- Context helps monotonically as DAMAGE CONTROL: none < news < macro < macro&news, in both settings. macro&news nearly halved the discretionary loss (-51.6% → -32.9%).
- Macro dominates; news is a small add-on (regime/breadth does most of the work).
- Best combo: macro&news + balanced (-0.44% vs strict, Sharpe ~3.0) — 'mostly follow the model, with full market context as a sanity check.'

Cross-check — Haiku 4.5, short window (January, ~19 days)

- Weak model + short window = no clear signal: deltas were sub-1pp (within run-to-run noise).
- macro HURT Haiku slightly but HELPED Sonnet → a source's value is model-dependent; you cannot evaluate sources on a weak model.
- Autonomy's damage only became obvious over the longer window — short horizons hide it.

Bottom line

- Trust the model; use context as a guardrail, not a license to deviate. Macro > news, and a capable model (Sonnet> Haiku) is required to use context at all.
- Caveat: one strong-momentum bull window where model_v4 is near-optimal— so 'follow the model' winning is expected. Decisive next test: a choppy/selloff OOS window.